

Nick Montana

Professor Yan Yan

Intro to Machine Learning

3/20/2018

Project 1 Report

Introduction: Project 1 consists of completing various functions. Significant code snippets and outputs can be found below:

computeCost.m:

```
% ===== YOUR CODE HERE =====
% Instructions: Compute the cost of a particular choice of theta
%               You should set J to the cost.

h = X*theta;      %This gives the h_theta values in a 97-length array.

sum = 0;

cost_sum = 0;
for i = 1:m
    a = h(i) - y(i);
    cost_sum = cost_sum + a^2;
end

J = cost_sum/(2*m);

% =====
```

computeCostMulti.m:

```
% ===== YOUR CODE HERE =====
% Instructions: Compute the cost of a particular choice of theta
%               You should set J to the cost.

h = X*theta;      %This gives the h_theta values in a 97-length array.

sum = 0;

cost_sum = 0;
for i = 1:m
    a = h(i) - y(i);
    cost_sum = cost_sum + a^2;
end

J = cost_sum/(2*m);

% =====
```

featureNormalize.m:

```
% ===== YOUR CODE HERE =====
% Instructions: First, for each feature dimension, compute the mean
%               of the feature and subtract it from the dataset,
%               storing the mean value in mu. Next, compute the
%               standard deviation of each feature and divide
%               each feature by it's standard deviation, storing
%               the standard deviation in sigma.
%
%               Note that X is a matrix where each column is a
%               feature and each row is an example. You need
%               to perform the normalization separately for
%               each feature.
%
% Hint: You might find the 'mean' and 'std' functions useful.

m = size (X, 1);
n = size (X, 2);

mu = mean (X, 1);
sigma = std (X, 1);

X_norm = X;

for j = 1:n
    for i = 1:m
        X_norm (i, j) = X (i, j) - mu (j);
    end
end

for j = 1:n
    for i = 1:m
        X_norm (i, j) = X_norm (i, j) / sigma (j);
    end
end

% =====
```

gradientDescent.m:

```
% ===== YOUR CODE HERE =====
% Instructions: Perform a single gradient step on the parameter vector
%               theta.
%
% Hint: While debugging, it can be useful to print out the values
%       of the cost function (computeCost) and gradient here.
%
%
h = X*theta;      %This gives the h_theta values in a 97-length array.
                  %This is basically the prediction in our model.

cost_sum = 0;     %Starting with sum being 0.

%%%%%%%%%%%%This updates theta_0 part%%%%%%%%%%%%
```

```

for i = 1:m          %This calculates the summation part
    a = h(i) - y (i);
    cost_sum = cost_sum + a*(X(i, 1));
end

update = (alpha / m) * cost_sum;    %This is the - " " part

temp0 = theta (1) - update;          %Updates theta and stores in temp
                                      % variable

%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
%Follows the same process from theta_0
cost_sum = 0;

for i = 1:m
    a = h(i) - y (i);
    cost_sum = cost_sum + a*(X(i, 2));
end

update = (alpha / m) * cost_sum;

temp1 = theta (2) - update;

theta (1) = temp0;                    %This actually updates theta
theta (2) = temp1;

end
% =====

```

gradientDescentMulti.m:

```

% ===== YOUR CODE HERE =====
% Instructions: Perform a single gradient step on the parameter vector
%               theta.
%
% Hint: While debugging, it can be useful to print out the values
%       of the cost function (computeCostMulti) and gradient here.
%

h = X*theta;          %This gives the h_theta values in a 97-length array.
                      %This is basically the prediction in our model.

temp = zeros (1, size (X, 2));

for j = 1:size (X, 2)

    cost_sum = 0;      %Starting with sum being 0.

    for i = 1:m        %This calculates the summation part
        a = h(i) - y (i);
        cost_sum = cost_sum + a*(X(i, j));
    end
end

```

```

        update = (alpha / m) * cost_sum;    %This is the - " " part

        temp (j) = theta (j) - update;      %Updates theta and stores in
                                           % temp variable
    end

    for i = 1:size (X, 2)
        theta (i) = temp (i);
    end

    % =====

```

normalEqn.m:

```

% ===== YOUR CODE HERE =====
% Instructions: Complete the code to compute the closed form solution
%               to linear regression and put the result in theta.
%
% -----

theta = inv(X'*X)*(X'*y);

% =====

```

plotData.m:

```

% ===== YOUR CODE HERE =====

plot (x,y, 'rx', 'MarkerSize', 10);      % Plot the data
ylabel('Profit in $10,000's');           % Set the y-axis label
xlabel('Population of City in 10,000s');  % Set the x-axis label

% Instructions: Plot the training data into a figure using the
%               "figure" and "plot" commands. Set the axes labels using
%               the "xlabel" and "ylabel" commands. Assume the
%               population and revenue data have been passed in
%               as the x and y arguments of this function.
%
% Hint: You can use the 'rx' option with plot to have the markers
%       appear as red crosses. Furthermore, you can make the
%       markers larger by using plot(..., 'rx', 'MarkerSize', 10)

% =====

```

warmUpExercise.m:

```

% ===== YOUR CODE HERE =====
% Instructions: Return the 5x5 identity matrix
%               In octave, we return values by defining which variables
%               represent the return values (at the top of the file)

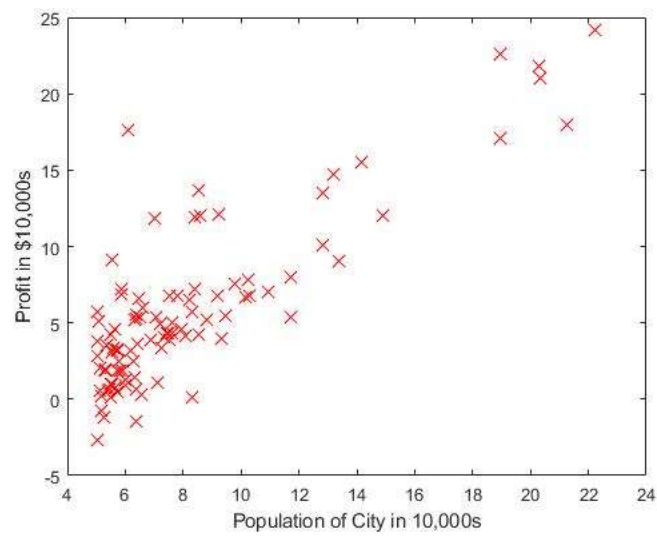
```

```
% and then set them accordingly.
```


```
A = eye (5);
```

```
% =====
```


ex1 Figure1:




ex1 Figure 1 with Line:



ex1 Figure 2:



ex1 Figure 3:



ex1 Output:

Command Window

New to MATLAB? See resources for [Getting Started](#).

Running warmUpExercise ...

5x5 Identity Matrix:

ans =

1	0	0	0	0
0	1	0	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	1

Program paused. Press enter to continue.

Plotting Data ...

Program paused. Press enter to continue.

Testing the cost function ...

With theta = [0 ; 0]

Cost computed = 32.072734

Expected cost value (approx) 32.07

With theta = [-1 ; 2]

Cost computed = 54.242455

Expected cost value (approx) 54.24

Program paused. Press enter to continue.

Running Gradient Descent ...

Theta found by gradient descent:

-3.630291

1.166362

Expected theta values (approx)

-3.6303

1.1664

For population = 35,000, we predict a profit of 4519.767868


For population = 70,000, we predict a profit of 45342.450129

Program paused. Press enter to continue.

Visualizing J(theta_0, theta_1) ...

f >> |

ex1_multi Figure 1:



ex1_multi Output:

```

Command Window
New to MATLAB? See resources for Getting Started.

Loading data ...
First 10 examples from the dataset:
x = [2104 3], y = 399900
x = [1600 3], y = 329900
x = [2400 3], y = 369000
x = [1416 2], y = 232000
x = [3000 4], y = 539900
x = [1985 4], y = 299900
x = [1534 3], y = 314900
x = [1427 3], y = 198999
x = [1380 3], y = 212000
x = [1494 3], y = 242500
Program paused. Press enter to continue.
Normalizing Features ...
Running gradient descent ...
Theta computed from gradient descent:
334302.063993
99411.449474
3267.012854

Predicted price of a 1650 sq-ft, 3 br house (using gradient descent):
$289221.547371
Program paused. Press enter to continue.
Solving with normal equations...
Theta computed from the normal equations:
89597.909543
139.210674
-8738.019112

Predicted price of a 1650 sq-ft, 3 br house (using normal equations):
$293081.464335
fx >> |

```